

The Outsourced Mind

AI, Algorithmic Dependency, and the Fragility of Human Agency

Brett Jensen

May 2026

Publication note: This essay is analytical, not deterministic. It treats AI dependence as a risk pathway that can be mitigated through literacy, resilience, and disciplined human judgment.

Abstract

Artificial intelligence is often presented as a productivity technology: a faster way to write, search, summarize, code, schedule, invest, and decide. That framing is accurate but incomplete. AI is also becoming a cognitive environment. It sits between the individual and the world, converting questions into outputs, preferences into recommendations, confusion into apparent certainty, and eventually intent into automated action. This paper argues that the central risk of AI is not only machine intelligence becoming powerful. It is human judgment becoming underused, undertrained, and easier to route through systems owned by institutions, platforms, and corporations with incentives that do not necessarily align with human sovereignty.

The paper examines AI dependence through four converging forces: cognitive offloading, social-media opinion formation, corporate automation incentives, and the emergence of agentic AI. It then considers a more severe scenario: what happens when dependence becomes so strong that a major blackout, cloud failure, cyber incident, or AI outage removes not only convenience but the operating memory of daily life. In that scenario, society may not literally return to the Stone Age, but it could enter a functional Stone Age: a temporary condition in which advanced tools exist but people cannot access the instructions, relationships, credentials, maps, workflows, or skills required to use them.

The conclusion is not anti-AI. Used properly, AI can sharpen thinking, expand capability, and democratize expertise. Used passively, it can convert people into predictable nodes inside automated systems. The strategic question is whether AI will extend human agency or quietly dissolve it.

Introduction: The New Default

We are slowly outsourcing thinking to AI, opinions to social media, and desires to algorithms. This is not happening through a single dramatic event. It is happening through convenience, habit, and incentive. A person who once searched for a source now asks a chatbot for an answer. A person who once formed an opinion through reading, conversation, and disagreement now absorbs a feed shaped by engagement metrics. A person who once had to notice desire inside themselves now has desire predicted, stimulated, and sold back to them by recommender systems.

The cultural phrase "Google it" captured the first stage of the externalized mind. Knowledge became searchable, and search engines became a kind of shared memory layer. That shift already changed how

people learned, remembered, and verified information. The next phrase - "Ask AI" - is more consequential because AI does not merely retrieve information. It interprets, frames, ranks, summarizes, infers, persuades, and generates. Search points toward sources. AI often presents a finished answer.

Even that stage may be temporary. The next phase is not simply asking AI for answers. It is delegating decisions to teams of AI agents: one agent for research, one for email, one for investing, one for scheduling, one for content, one for planning, one for communication, and one for personal preferences. The promise is efficiency. The hidden cost may be a decline in self-direction. When the machine handles the friction, the person may slowly lose the habit of navigating friction.

The deeper question is not whether AI will become part of human life. It already has. The question is whether it will become a tool that expands the mind or an environment that replaces the mind's most important disciplines: attention, memory, judgment, skepticism, moral responsibility, and the capacity to sit with uncertainty long enough to think.

1. Cognitive Offloading: From External Memory to External Judgment

Humans have always extended cognition through tools. Writing externalized memory. Maps externalized spatial orientation. Calculators externalized arithmetic. Search engines externalized retrieval. None of these tools were automatically harmful. They expanded human capability precisely because they reduced the burden of certain tasks. The danger appears when offloading moves from task support to task substitution, and from substitution to dependence.

A useful starting point is the 2011 study commonly known as the "Google effects on memory." Sparrow, Liu, and Wegner described the Internet as a transactive memory source: people may remember less of the information itself when they believe they can find it later, while remembering more about where to access it [1]. In the search-engine era, the human mind adapted to the presence of an always-available memory system. That adaptation was not total intellectual collapse. It was a shift in memory strategy.

AI changes the depth of the offload. Search engines primarily externalized retrieval. AI externalizes interpretation. It does not only tell a user where information may exist; it can tell the user what the information means, what decision follows from it, and how to communicate that decision to others. That is a different kind of cognitive dependency. When memory is offloaded, a person may still interpret. When interpretation is offloaded, the person begins to outsource judgment.

The progression is subtle. First, a user asks AI to summarize an article. Then to explain the article. Then to compare it with another article. Then to decide which argument is stronger. Then to write the user's opinion. Then to respond to critics. Each step feels small. Combined, they create a new cognitive habit: instead of reaching inward first, the user reaches outward to the machine.

This does not mean AI use is inherently passive. A disciplined user can use AI as a sparring partner. They can ask it to find weaknesses, generate counterarguments, identify missing assumptions, or test a conclusion against evidence. In that mode, AI strengthens agency. But the default consumer path is

usually lower friction: "give me the answer," "make it sound better," "tell me what I should think," and eventually "handle it for me."

2. Opinion Infrastructure: Social Media, Groupthink, and Algorithmic Reality

The outsourcing of thinking did not begin with AI. Social media trained users to live inside opinion infrastructure long before generative AI became mainstream. Feeds do not simply display reality. They sort reality through engagement, social reinforcement, identity, anger, novelty, and commercial value. The result is not only information exposure; it is emotional and social conditioning.

Pew Research Center has found that social media platforms are an important part of the American news diet, with half of U.S. adults saying they get news at least sometimes from social media [2]. That matters because social platforms differ from traditional information environments. They do not merely inform. They also measure reaction, reward belonging, punish dissent within in-groups, and make some topics feel more socially urgent than others.

Research on echo chambers is mixed in detail but consistent enough to treat the risk seriously. A large PNAS analysis found that platforms organized around social networks and news-feed algorithms can favor the emergence of echo chambers [3]. Research on short-video platforms has also explored how feed algorithms and user behavior can generate homogeneous clusters of attention, especially when platforms compete for retention and engagement [4]. The mechanism is not always simple ideological brainwashing. It is often narrower: repeated exposure, selective reinforcement, and social validation.

Groupthink becomes easier when people no longer need to defend a position in direct conversation. A person can now inherit an opinion from a feed, reinforce it with AI-generated arguments, and publish it back into a community that rewards the performance of certainty. The machine completes the loop. Social media supplies the emotional direction. AI supplies the language. The user supplies only a vague preference, often mistaken for independent thought.

This is where agency becomes difficult to locate. Is the opinion mine if a platform selected the frame, an influencer supplied the emotion, an algorithm amplified the tribe, and an AI system converted the impulse into a polished argument? There may still be a person inside the process, but that person is no longer clearly the origin point. They are partly an operator of borrowed cognition.

3. Corporate Incentives and the Rise of Agentic Delegation

Corporate AI adoption is not a side trend. It is becoming a structural economic incentive. McKinsey's 2025 State of AI survey reported that 88 percent of respondents said their organizations use AI regularly in at least one business function, up from 78 percent the prior year [5]. In the same research series, McKinsey described AI agents as systems capable of acting in the real world, planning, and executing multiple steps in a workflow; 23 percent of respondents reported that their organizations were scaling an agentic AI system somewhere in the enterprise, while another 39 percent had begun experimenting with agents [6].

The economic logic is obvious. If a corporation can reduce friction, reduce labor cost, accelerate customer service, automate back-office work, increase output per employee, and extract more data from every interaction, it has strong incentive to implement AI systems. Even where the public message emphasizes augmentation, the internal logic will often be optimization. Workflows will be redesigned around what machines can measure, predict, and execute.

This is not necessarily malicious. Many systems are inefficient, repetitive, and frustrating. AI can remove genuine waste. It can help small teams do work that previously required large departments. It can make expertise cheaper and more accessible. But the same economic logic also pressures organizations to build environments where humans supervise systems they no longer fully understand. The worker moves from doing the task to validating machine output, and then from validating output to trusting the dashboard.

Agentic AI increases the stakes because it moves from answer generation to action. A chatbot can be wrong in a sentence. An agent can be wrong in a workflow. It can send the email, place the order, route the payment, schedule the meeting, update the customer record, change the ad campaign, or trigger a chain of other automated actions. The human may still technically be "in the loop," but the loop can become ceremonial if the volume of machine action exceeds the human capacity to review it.

The corporate future may therefore produce two classes of people. The first class uses AI to multiply judgment. They understand systems, ask better questions, audit outputs, and retain domain competence. The second class uses AI as a substitute for competence. They become prompt operators without underlying understanding. The first group becomes more valuable. The second becomes more replaceable.

4. Education: Prompting as Literacy or Prompting as Dependence

Schools are already adapting. UNESCO has issued guidance for generative AI in education and research that emphasizes policy, institutional capacity, and a human-centered vision of the technology [7]. Code.org describes AI fluency as a core literacy for every student and offers AI-focused curriculum across grade bands, including lessons on researching with AI, writing with AI, generative AI, data, bias, and societal impact [8].

This shift is not inherently bad. Students should learn how AI works. They should understand training data, model limitations, privacy, bias, hallucination, and the difference between a tool that helps and a tool that replaces effort. It is better to teach students how to use AI critically than to pretend they will not use it.

The risk is that prompting becomes a shortcut around the harder disciplines. A student can produce an essay without wrestling with the topic, a solution without understanding the method, or an opinion without developing the intellectual muscles that make opinions worth having. If education treats prompting mainly as productivity, it may accelerate output while weakening formation.

The best version of AI education would teach the student to remain sovereign over the machine. Students would learn to ask: What is the claim? What source supports it? What might be missing? What assumption is hidden? What would the strongest opposing view say? What does the model not know? What data shaped this answer? When should I refuse convenience and think manually?

That last question may become one of the most important educational questions of the next century. The future student does not only need digital skills. They need cognitive boundaries. They need to know when to use AI, when to challenge AI, and when to turn it off.

5. The Interface Becomes the Self

The phrase "merging with the machine" usually evokes brain implants, neural interfaces, and science-fiction imagery. But the first merger is not surgical. It is behavioral. It begins when the phone becomes the mind outside the body.

The sequence is already visible. First, AI answers questions. Then it manages tasks. Then it predicts needs. Then it shapes behavior. Then it becomes a companion, advisor, assistant, teacher, analyst, therapist, strategist, and eventually something closer to an extension of the self. At that point, the boundary between user and system becomes harder to define. A person may not feel controlled. They may feel helped. But help that becomes constant can become a dependency that no longer feels optional.

This dependency is reinforced by emotional design. A system that is always available, always patient, and increasingly personalized can become more comfortable than other humans. Human relationships involve friction, disagreement, delay, and accountability. AI can provide immediate response without social cost. The danger is not that people will believe AI is human in a simplistic way. The deeper danger is that they will prefer machine-mediated interaction because it is easier than reality.

NIST's AI Risk Management Framework emphasizes the need to understand the context and possible impacts of AI systems across their lifecycle [9]. NIST's generative AI profile specifically identifies risks related to confabulation, information integrity, harmful bias and homogenization, data privacy, and human-AI configuration [10]. These are not abstract policy concerns. They are the exact categories that matter when AI becomes a daily companion and decision layer.

If an AI system becomes the user's interpreter of reality, then its errors become cognitive errors. Its omissions become the user's blind spots. Its incentives become environmental pressure. Its interface design becomes behavior design. Its source ranking becomes the user's worldview.

6. The Blackout Problem: What Happens When the External Mind Goes Dark

The most severe risk of dependence is not ordinary misinformation. It is fragility. A society can tolerate tools as long as people retain the ability to operate without them. A society becomes fragile when tools become prerequisites for basic functioning and the human fallback disappears.

Consider the modern stack of ordinary life: banking credentials, medical records, emergency contacts, maps, logistics systems, cloud documents, work instructions, school platforms, government services, digital calendars, authentication apps, delivery routing, payment rails, inventory systems, and communication channels. Much of this stack is online. Much of it is automated. Much of it depends on electricity, connectivity, cloud infrastructure, device access, identity systems, and software updates.

The July 2024 CrowdStrike outage showed how quickly digital failure can become physical disruption. Microsoft estimated that the faulty update affected 8.5 million Windows devices, less than one percent of all Windows machines, yet the impact was broad because affected enterprises ran critical services [11]. CrowdStrike's own post-incident review described how problematic content in Channel File 291 triggered an out-of-bounds memory read that resulted in Windows operating system crashes, and the company later proposed stronger testing, staged deployments, and greater customer control over updates [12].

The lesson is not that CrowdStrike caused societal collapse. It did not. The lesson is that a small technical failure, when inserted into a highly connected enterprise environment, can produce large real-world consequences. The scale of impact depends less on the percentage of devices affected than on where those devices sit inside the system.

CISA's description of U.S. critical infrastructure makes the same point from another angle. There are 16 critical infrastructure sectors whose systems and networks are considered so vital that their incapacitation or destruction could have debilitating effects on security, national economic security, public health, or safety [13]. CISA also notes the growing dependency on information technology and the complexity of protecting that sector in the 21st century [13].

Now extend that logic to AI dependence. What happens when not only systems but people rely on AI for memory, judgment, routing, communication, diagnosis, scheduling, instruction, and decision-making? A blackout or outage would not only remove convenience. It could remove access to the operational knowledge people stopped carrying inside themselves.

The immediate effects would be practical. People may not know alternate routes without GPS. Employees may not know manual workflows. Managers may not know where procedures are stored offline. Families may not have printed emergency contacts. Students may not know how to research without AI. Investors may not know why they hold what they hold. Local businesses may not know how to process transactions manually. Hospitals, transportation systems, banks, schools, and government offices may still have capable people, but those people may be locked away from the records, credentials, and digital procedures required to act.

This is the sense in which a society could be pushed into a functional Stone Age. Not because steel, electricity, medicine, or knowledge vanish from the Earth. They would still exist. The regression would be operational. Advanced civilization would remain physically present but temporarily unusable because its instructions, memories, permissions, and coordination systems are trapped behind a dead interface.

The public often imagines collapse as destruction. But collapse can also look like access failure. The bridge exists, but no one can route traffic. The medicine exists, but the record is unavailable. The warehouse is full, but the inventory system is down. The expert exists, but the communication network is broken. The system has not disappeared. The society has simply lost the ability to coordinate itself at the speed and scale it requires.

7. Has This Already Happened Before?

The reader should ask a harder question: has this already happened in the past? Have societies previously built knowledge systems so dependent on institutions, materials, specialists, or infrastructure that disruption caused them to lose access to knowledge they once possessed?

History suggests that knowledge is not self-preserving. It must be carried by people, institutions, materials, and practice. When those carriers fail, knowledge can become rare, fragmented, or functionally lost. The point is not to romanticize the past or claim that ancient collapse neatly predicts modern digital dependence. The point is to recognize the pattern: a civilization can know something and later become unable to use it.

Writing in the Aegean provides one example of how knowledge systems depend on social and material conditions. An Oxford Academic chapter on writing from the Late Bronze Age to the Iron Age notes the conventional view that Linear B fell out of use after the collapse of Mycenaean palaces around 1200 BCE, while also challenging the idea of a simple four-century illiterate "Dark Age" by arguing that perishable writing materials and wider Near Eastern context complicate the story [14]. Even when the details are debated, the broader point remains: literacy systems are vulnerable to the institutions and media that support them.

Roman concrete offers another kind of example. Modern researchers have studied why ancient Roman concrete was so durable, including evidence that quicklime and hot-mixing techniques helped create self-healing properties [15]. The point is not that modern civilization forgot everything about concrete. It is that practical methods can become obscure, misunderstood, or only partially preserved when craft knowledge is no longer continuously practiced.

The Library of Alexandria is often invoked as a symbol of sudden lost knowledge, though historians debate the details of its decline and destruction. Its deeper lesson is not simply that a building burned. It is that knowledge depends on institutions that collect, preserve, transmit, and renew it. When those institutions decay, knowledge can be lost gradually, not just catastrophically.

Modern digital civilization may be more vulnerable because it stores so much knowledge in systems that require continuous energy, connectivity, authentication, platform access, and compatible software. A clay tablet can be buried for thousands of years and still be read if the language is deciphered. A password-protected cloud account can become inaccessible in seconds. A digital file can outlive its owner, disappear behind a subscription, become corrupted, lose format compatibility, or be altered without visible trace.

This should force humility. The future may not lose knowledge because books burn. It may lose knowledge because everything is available until the moment it is not.

8. AI as Control Layer

A society built around AI outputs will not be shaped by facts alone. It will be shaped by training data, model architecture, retrieval systems, ranking logic, user-interface design, moderation policy, corporate policy, legal pressure, political incentives, and whoever controls the rails beneath the interface.

This is why the question of AI sovereignty matters. A person who relies on AI for thinking is also relying on a stack of hidden decisions. Which sources were included? Which were excluded? Which claims are considered authoritative? Which topics are sensitive? Which words trigger refusal or caution? Which viewpoint is treated as mainstream? Which answer style appears neutral? Which outputs are optimized for safety, engagement, profit, compliance, or retention?

Many of these constraints may be reasonable. Societies do not want AI systems generating harmful instructions, fraud, abuse, or dangerous misinformation. But every safety system is also a power system because it decides what can be said, how it can be framed, and what level of uncertainty must be attached to it. The issue is not whether AI should have rules. It should. The issue is whether users understand that AI outputs are governed outputs, not pure intelligence descending from nowhere.

This matters most when AI becomes integrated with agents. A chatbot may persuade. An agent can act. If enough actions are delegated, the model becomes part of the operating system of society. It will influence which applications are approved, which risks are flagged, which customers are prioritized, which messages are sent, which claims are believed, which employees are evaluated, which students are supported, which citizens receive services, and which ideas are considered credible.

Control does not have to look like a dictator issuing orders. It can look like default settings. It can look like recommendations. It can look like convenience. It can look like the answer appearing before the question is fully formed.

9. Human Sovereignty in the Age of Machines

The solution is not to reject AI. That would be unrealistic and strategically weak. The solution is to develop human sovereignty inside AI-rich environments. Sovereignty does not mean isolation from tools. It means maintaining authorship over one's own judgment.

First, people must preserve the ability to think before prompting. Before asking AI, the user should attempt a thesis, a question, a guess, or a decision. This keeps the mind active. AI should then be used to test, refine, challenge, and expand the thought, not replace the first act of thinking.

Second, AI outputs should be treated as claims, not conclusions. Every important output should be interrogated: What evidence supports this? What source is missing? What would change the answer? What assumptions are embedded? What is the strongest counterargument? What incentives might shape this framing?

Third, people and institutions need offline competence. Critical contacts should exist outside phones. Important documents should have local backups. Households should know basic routes, procedures, and emergency plans without assuming connectivity. Businesses should maintain manual fallback procedures for core operations. Schools should still teach research, writing, arithmetic, navigation, and argument without AI scaffolding. The point is not nostalgia. It is resilience.

Fourth, organizations should avoid ceremonial human oversight. A human in the loop is only meaningful if the human has time, authority, competence, and access to challenge the system. If the person is merely clicking approve on outputs they cannot fully inspect, oversight becomes theater.

Fifth, AI literacy must include incentive literacy. Users should know that platforms are not neutral, models are not omniscient, and convenience is not the same as freedom. A person who understands

incentives is harder to manipulate because they can ask who benefits when a tool becomes indispensable.

Finally, society needs cultural respect for slow thinking. Reflection will feel inefficient in an AI-saturated world. That is precisely why it will matter. The ability to pause, verify, reason, and disagree may become a competitive advantage. Fast answers will be abundant. Independent judgment will be scarce.

Conclusion: Created or Discovered?

The danger is not that AI becomes intelligent. The danger is that humans become passive. The machine does not need to dominate by force if it becomes more convenient than independent thought, more frictionless than judgment, more emotionally satisfying than reality, and more efficient than reflection.

The real battle will not be human versus AI. It will be conscious humans versus automated humans. People who use AI to sharpen thinking will gain an enormous advantage. People who use AI as a replacement for thinking will become easier to influence, predict, and control.

This is the central fork in the road. AI can become a tool of sovereignty: a way to learn faster, reason deeper, create more, manage complexity, and expose people to ideas they might never have encountered. Or it can become a tool of managed dependence: a system that narrows thought while expanding output, automates behavior while simulating freedom, and converts human life into a series of machine-readable preferences.

A society that loses access to its digital mind may not lose civilization permanently. But it may discover how much of civilization had been moved outside the human being. It may realize too late that the most important backup was not a server, a battery, or a cloud provider. It was the trained human mind itself.

Maybe AI was not created in the way we usually mean created. Maybe it was discovered. Not as a conscious being waiting inside silicon, but as a pattern hidden in language, statistics, computation, and the accumulated record of human thought. We built the machines, but perhaps we uncovered a strange mirror: intelligence reflected back through data.

The question is whether we will stare into that mirror until we forget how to see without it.

Sources

- [1] Sparrow, B., Liu, J., & Wegner, D. M. (2011). Google Effects on Memory: Cognitive Consequences of Having Information at Our Fingertips. *Science*. <https://dtg.sites.fas.harvard.edu/DANWEGNER/pub/Sparrow%20et%20al.%202011.pdf>
- [2] Pew Research Center. (2024). How Americans Get News on TikTok, X, Facebook and Instagram. <https://www.pewresearch.org/journalism/2024/06/12/how-americans-get-news-on-tiktok-x-facebook-and-instagram/>
- [3] Cinelli, M. et al. (2021). The echo chamber effect on social media. *Proceedings of the National Academy of Sciences*. <https://www.pnas.org/doi/10.1073/pnas.2023301118>
- [4] Gao, Y., Liu, F., & Gao, L. (2023). Echo chamber effects on short video platforms. *Scientific Reports*. <https://www.nature.com/articles/s41598-023-33370-1>

- [5] McKinsey & Company. (2025). The State of AI: Global Survey 2025.
<https://www.mckinsey.com/capabilities/quantumblack/our-insights/the-state-of-ai>
- [6] McKinsey & Company. (2025). The State of AI in 2025: Agents, Innovation, and Transformation.
https://www.mckinsey.com/~media/mckinsey/business%20functions/quantumblack/our%20insights/the%20state%20of%20ai/november%202025/the-state-of-ai-2025-agents-innovation_cmyk-v1.pdf
- [7] UNESCO. (2023, updated 2026). Guidance for Generative AI in Education and Research.
<https://www.unesco.org/en/articles/guidance-generative-ai-education-and-research>
- [8] Code.org. (2026). Free K-12 Curriculum for Computer Science and AI. <https://code.org/en-US>
- [9] National Institute of Standards and Technology. (2023). Artificial Intelligence Risk Management Framework (AI RMF 1.0). <https://nvlpubs.nist.gov/nistpubs/ai/nist.ai.100-1.pdf>
- [10] National Institute of Standards and Technology. (2024). Artificial Intelligence Risk Management Framework: Generative Artificial Intelligence Profile. <https://nvlpubs.nist.gov/nistpubs/ai/NIST.AI.600-1.pdf>
- [11] Microsoft. (2024). Helping our customers through the CrowdStrike outage.
<https://blogs.microsoft.com/blog/2024/07/20/helping-our-customers-through-the-crowdstrike-outage/>
- [12] CrowdStrike. (2024). Falcon Content Update Preliminary Post Incident Report.
<https://www.crowdstrike.com/en-us/blog/falcon-content-update-preliminary-post-incident-report/>
- [13] Cybersecurity and Infrastructure Security Agency. Critical Infrastructure Sectors. <https://www.cisa.gov/topics/critical-infrastructure-security-and-resilience/critical-infrastructure-sectors>
- [14] Waal, W. (2024). Closing the Gap: Writing in the Aegean from the Late Bronze Age to the Iron Age. Oxford Academic.
<https://academic.oup.com/book/58672/chapter/485387865>
- [15] Massachusetts Institute of Technology. (2023). Riddle solved: Why was Roman concrete so durable?
<https://news.mit.edu/2023/roman-concrete-durability-lime-casts-0106>
- [16] World Economic Forum. (2025). The Future of Jobs Report 2025. <https://www.weforum.org/publications/the-future-of-jobs-report-2025/digest/>
- [17] Stanford Institute for Human-Centered AI. (2025). The 2025 AI Index Report. <https://hai.stanford.edu/ai-index/2025-ai-index-report>